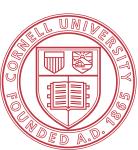

CS5112: Algorithms and Data Structures for Applications

Lecture 22: Computability; Max-flow for vision

Ramin Zabih

Some content from: Wikipedia/Google image search;
Olga Veksler & Yuri Boykov



Cornell University



**CORNELL
TECH**

Administrivia

- Wednesday lecture: review session
 - Come with questions!
- Monday: final exam
 - Students will be assigned to the auditorium or to 071
 - We will let you know by email which room you are in
- The following will not be on the exam
 - Previous lecture (SGD)
 - Second half of today's lecture (max flow)
 - Lower bounds and approximation algorithms

Lecture Outline

- Bad news: Some problems turn out to be impossible to solve
 - Uncomputable problems
 - Ex: determining if a hash function has collisions
- Good news: Some problems are quite solvable
 - Binary image denoising
 - Max flow/min cut (“Graph cuts”)
 - Applications and extensions

The bad news

- Some problems are impossible to solve
- You need to be able to recognize them

Uncomputability

- Some well-defined problems are simply impossible
 - No matter what your hardware, they are uncomputable
- You need to be able to recognize them
 - Similar to NP-hardness
 - BUT: NP-hard problems can be solved
 - just not very fast

Halting problem

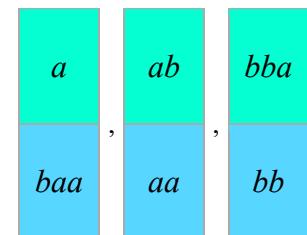
- Most famous example: the halting problem
 - Does a particular program run forever on a specific input?
- This seems easy to avoid, but...
- Consider a new problem such that, if you could solve it, would allow you to solve the halting problem
- There are some very subtle example of this
 - Just like with NP-hardness

When is a problem uncomputable?

- This is very difficult, but there are some rules and clues
- Any non-trivial property of a program is uncomputable
 - Rice's theorem
- Anything you can solve with exhaustive search is computable
 - Infinite/unbounded input is a key clue!
- Small differences in problem formulation can change a computable problem into an uncomputable one

A children's game

- We are given blocks with symbols such as a,b,c. Each block has a top and a bottom.
- There are certain types of blocks, such as a block with “ab” above “bc”.
- We have as many blocks of each type as we want.
- Can we find a series of blocks so that the top and bottom symbols are the same?



Game examples

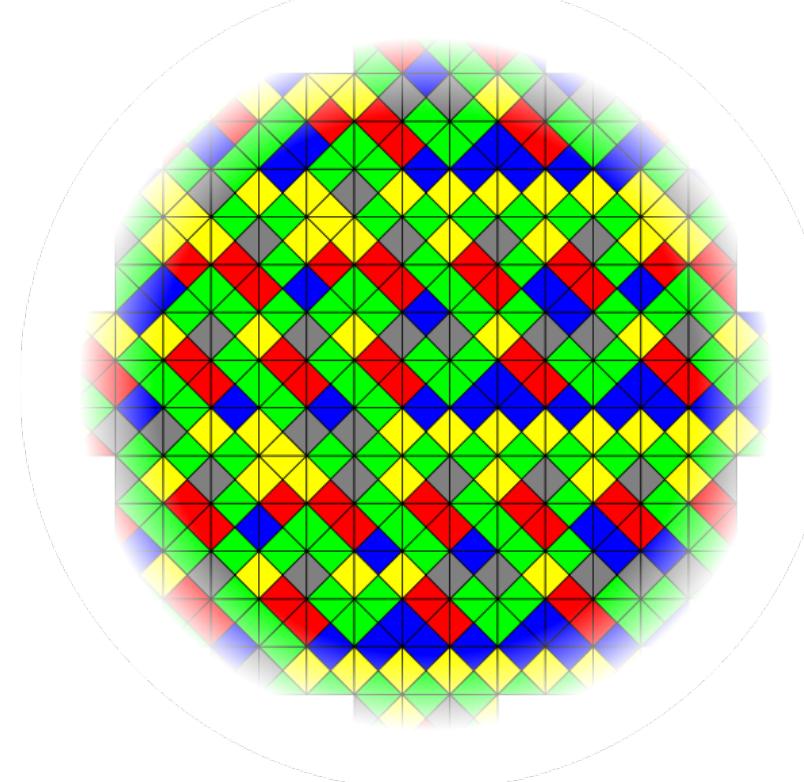
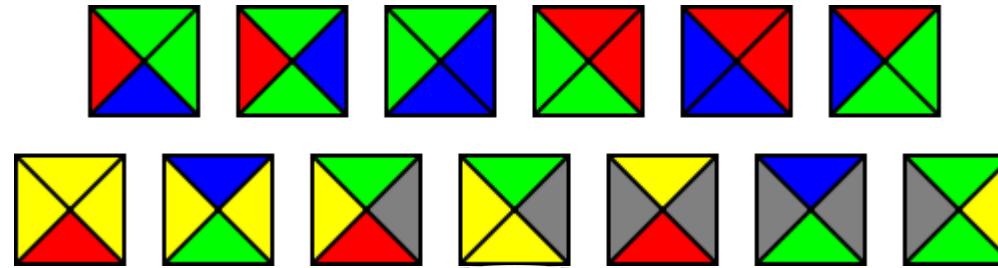
- Example 1:

a	ab	bba
baa	aa	bb
$i = 1$	$i = 2$	$i = 3$

Can we solve the children's game?

- For a binary alphabet (example 1) this is decidable
- For an alphabet with 7 characters or more?
 - undecidable!
- For an alphabet with 3 characters (example 2)?
 - open question
- What if we can use no more than k blocks, including copies, in our solution?
 - It's decidable but NP-hard

Tiling the plane



[Source](#)



CORNELL
TECH

The good news

- Some challenging looking problems are solvable
- Let's start by showing off a cool vision application

Interactive Digital Photomontage

**Aseem Agarwala, Mira Dontcheva,
Maneesh Agrawala, Steven
Drucker, Alex Colburn,
Brian Curless, David Salesin,
Michael Cohen**

University of Washington & Microsoft Research







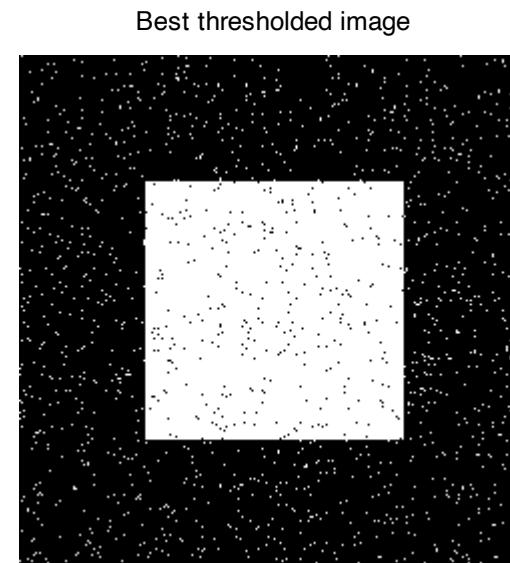
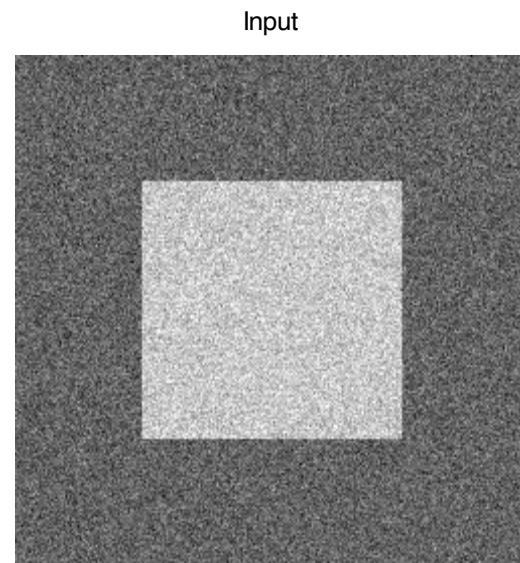






Image denoising

- Suppose we want to find a bright object against a dark background
 - But some of the pixel values are slightly wrong



Optimization viewpoint

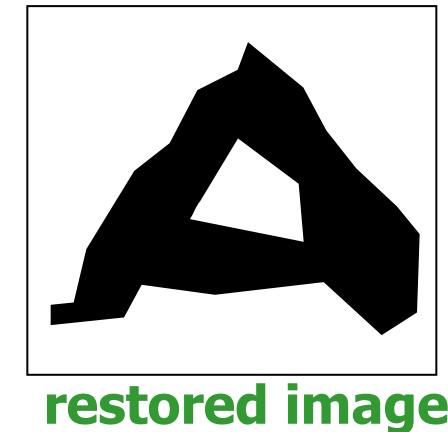
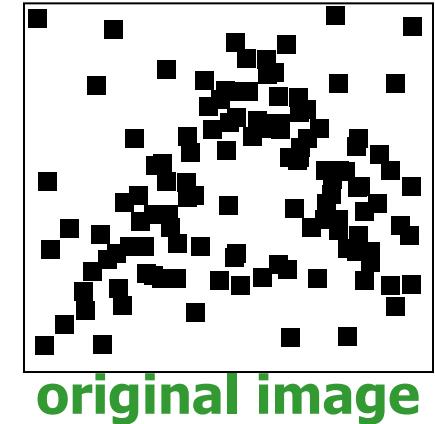
- Find best (least expensive) binary image
 - Costs: C1 (labeling) and C2 (boundary)
- C1: Labeling a dark pixel as foreground
 - Or, a bright pixel as background
- If we only had labeling costs, the cheapest solution is the thresholded output
- C2: The length of the boundary between foreground and background
 - Penalizes isolated pixels or ragged boundaries

Binary labeling problems

- There is a natural multi-label version
 - But 2 labels are surprisingly important
 - Lots of nice applications
 - Same basic ideas used for more labels
- Fast exact solution!
 - Turn a hard problem into a problem you can solve (reduction)
 - Due to Hammer (1965) originally
 - Job scheduling problems: Stone (1977)
 - Images: Greig, Porteus & Seheult (1989)

Binary Labeling: Motivation

- Suppose we receive a noisy fax:
 - Some black pixels in the original image were flipped to white pixels, and some white pixels were flipped to black
- We want to try to clean up (or **restore**) the original image:
- This problem is called **image restoration (denoising)**

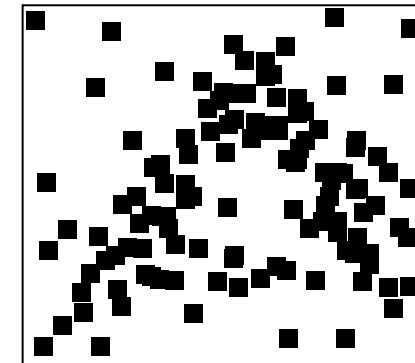


Binary Labeling Problem: Motivation

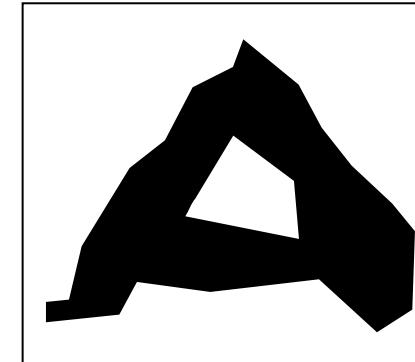
- Fax image is binary: each pixel is either
 - black (stored as 0)
 - or white (stored as 1)

■ What we know:

- 1) In the restored image, each pixels should also be either black (0) or white (1)
- 2) **Data Constraint**: if a pixel is black in the original image, it is more likely to be black in the restored image; and a white pixel in the original image is more likely to be white in the restored image
- 3) **Prior Constraint**: in the restored image, white pixels should form coherent groups, as should black pixels



original image



restored image

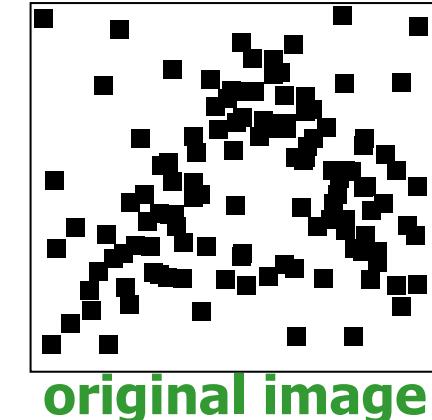


Binary Labeling Problem

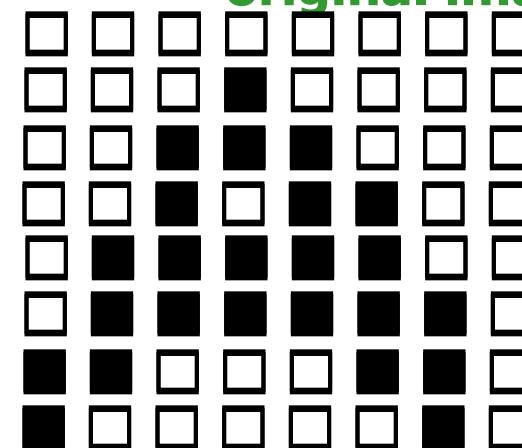
- We can formulate our restoration problem as a labeling problem:
 - Labels: black (0) and white (1)
 - Set of sites: all image pixels
 - Will use “sites” or pixels interchangeably

Assign a label to each site (either black or white) s.t.

- If a pixel is black (white) in the original image, it is more likely to get the black (white) label
- Black labeled pixels tend to group together, and white labeled pixels tend to group together



original image

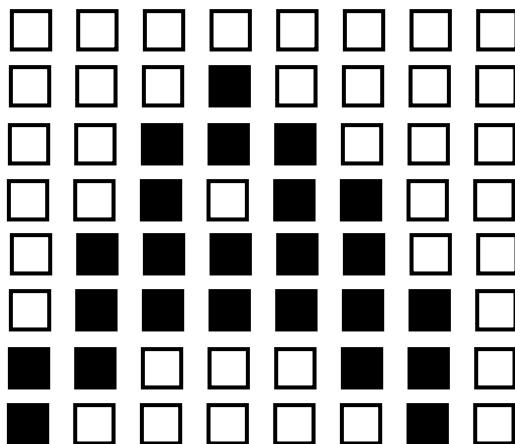


set of sites P , and
one possible labeling

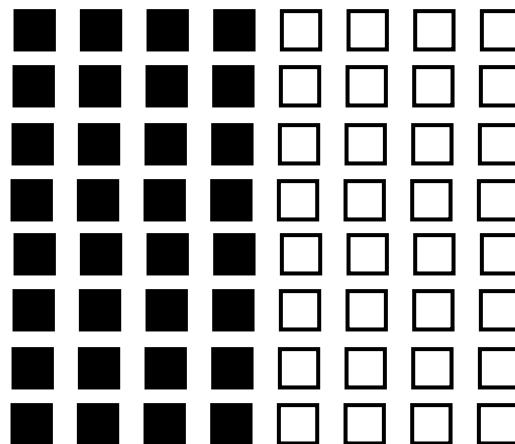


Binary Labeling Problem: Constraints

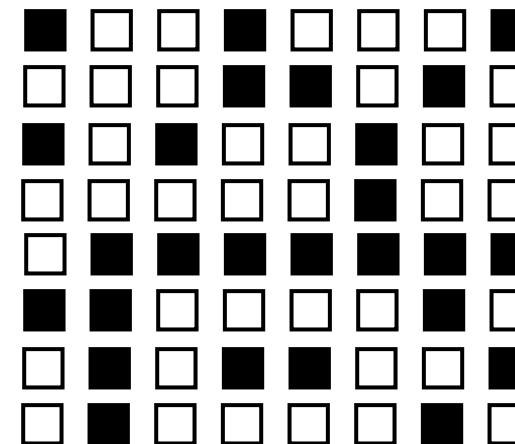
- If a pixel is black (white) in the original image, it is more likely to get the black (white) label
- Black labeled pixels tend to group together, and white labeled pixels tend to group together



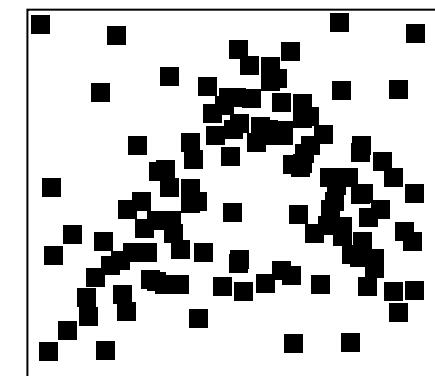
good labeling



bad labeling
(constraint 1)



bad labeling
(constraint 2)



original image



Binary Labeling Problem: Data Constraints

- How can we find a good labeling?
 - i.e., satisfying our constraints

- Formulate and minimize an energy function on labelings L :

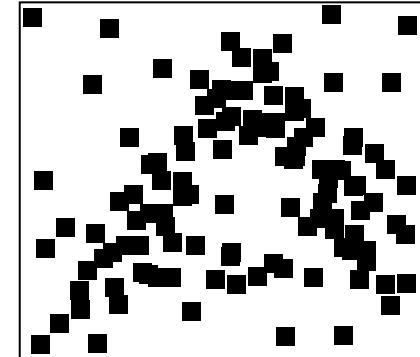
Let L_p be the label assigned to pixel p

- $L_p = 0$ or $L_p = 1$

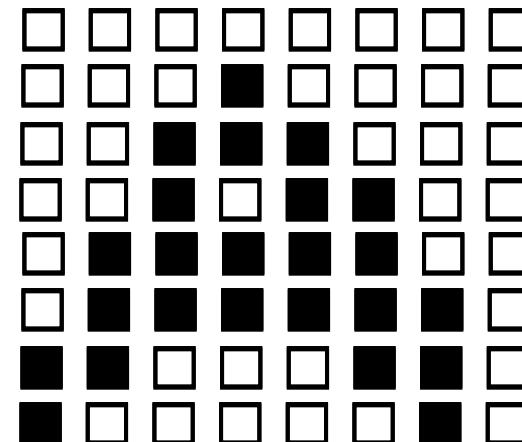
Let $I(p)$ be the intensity of pixel p in the original image

Data constraint is modeled by the Data Penalty $D_p(L_p)$

- $D_p(0) < D_p(1)$ if $I(p) = 0$
- $D_p(0) > D_p(1)$ if $I(p) = 1$



original image

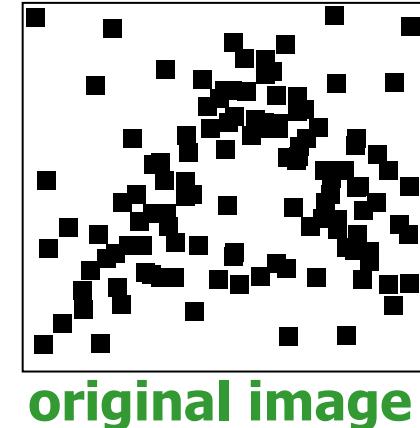


a good labeling L

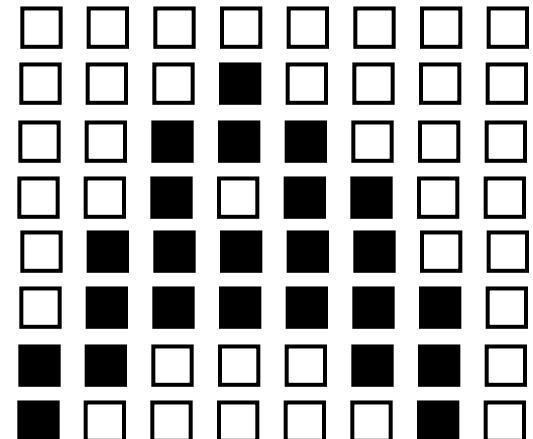


Binary Labeling Problem: Data Constraints

- In our example, we can make
 - if $I(p) = \text{black}$ then
 - $D_P(\text{black}) = 0$
 - $D_P(\text{white}) = 4$
 - if $I(p) = \text{white}$ then
 - $D_P(\text{black}) = 4$
 - $D_P(\text{white}) = 0$
- To figure out how well image as a whole satisfies the data constraint, sum up data terms over all image pixels: $\sum_p D_P(L_p)$



original image



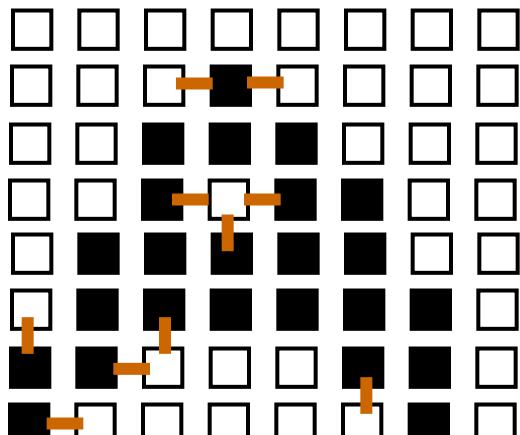
a good labeling L



Binary Labeling Problem: Prior Constraints

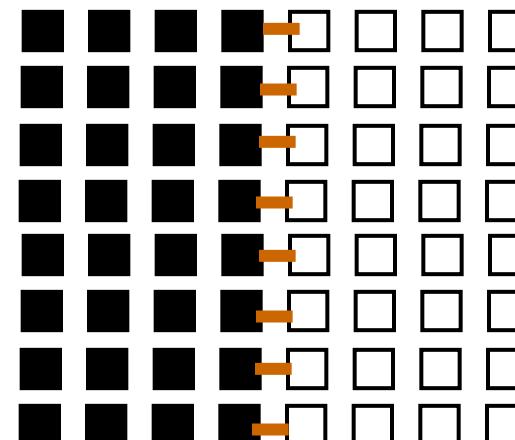
- To model our prior constraints, simply count the number “discontinuities” in a labeling L :
 - There is a discontinuity between pixels p and q if labels p and q have different labels
 - The more pixels with equal labels group together in coherent groups, the less discontinuities there are

**some discontinuities
are marked**



31 discontinuities

**all discontinuities
are marked**



8 discontinuities



Binary Labeling: Data Constraints

- How to count the number of discontinuities in a labeling L ?
 - Let N be the set of all adjacent pixels
 - Thus our N consists of edges between immediately adjacent pixels

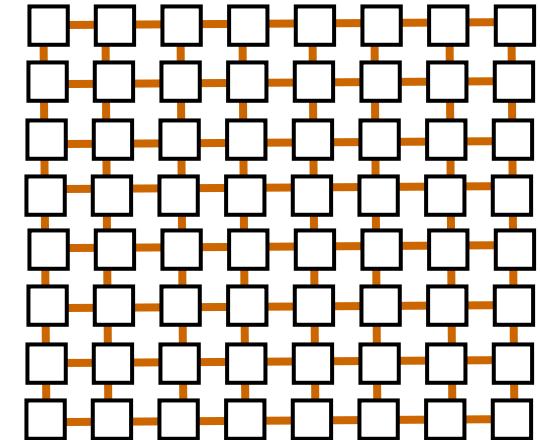
Let $I(b)$ be the identity function

$I(b) = 1$ when its argument b is true

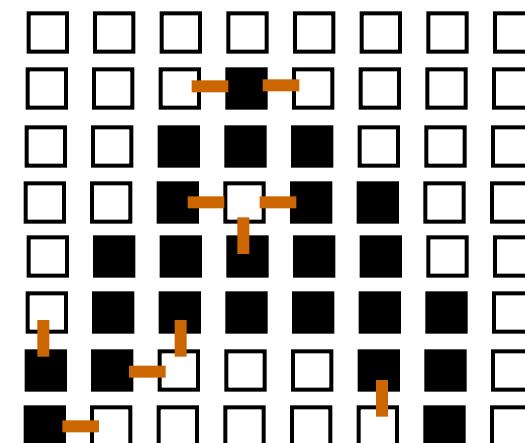
$I(b) = 0$ when its argument b is false

To count all discontinuities in a labeling L ,

$$\sum_{pq \in N} (L_p \neq L_q)$$



N consists of all edges



31 discontinuities



Binary Labeling Problem: Energy Formulation

- Our final energy, which takes into consideration the data and the prior constraints is:

$$E(L) = \sum_p D_P(L_p)$$

data term

$$+ \lambda \sum_{pq \in N} (L_p \neq L_q)$$

prior term

Data term says that in a good labeling L pixels should be labeled as close as possible to their colors in the original (noisy) image

Prior term says that in a good labeling L pixels should be labeled as to form spatially coherent blocks (as few discontinuities as possible)



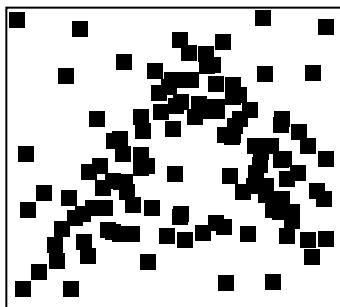
Binary Labeling Problem: Energy Formulation

$$E(L) = \sum_p D_P(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

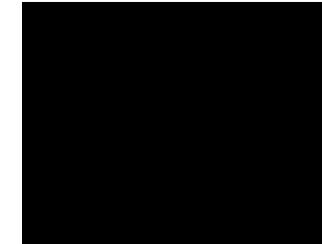
data term **prior term**

- Note that data term and prior term want different things:

Best labeling as far as the data term is concerned is the original image:



Best labeling considering only the prior term is completely black or completely white:

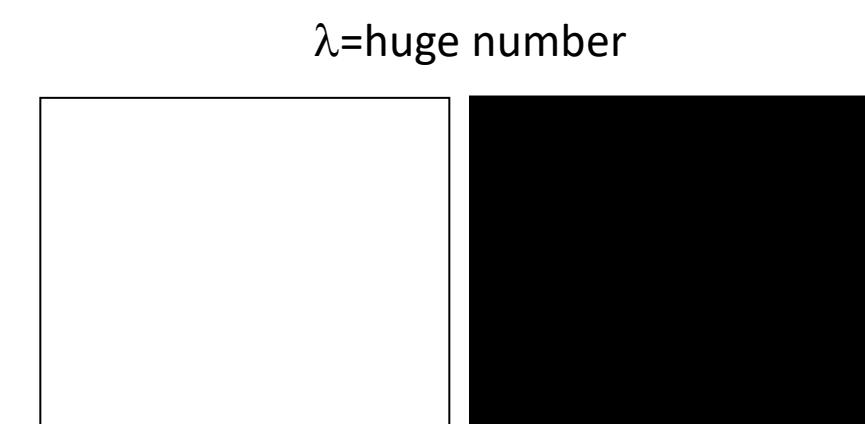
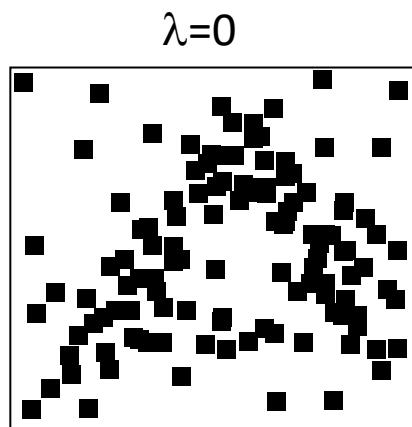


Binary Labeling Problem: Energy Formulation

$$E(L) = \sum_p D_P(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

data term **prior term**

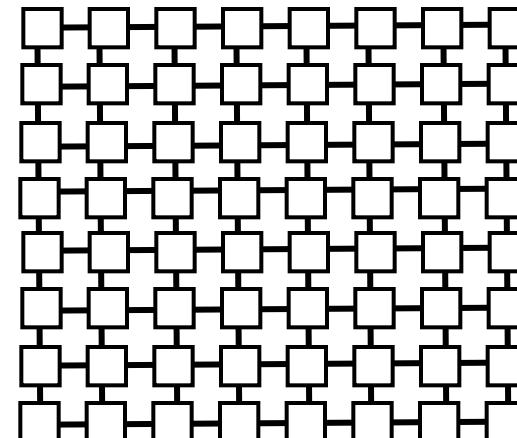
- λ serves as a balancing parameter between the data term and the prior term:
 - The larger the λ , the less discontinuities in the optimal labeling L
 - The smaller the λ , the more the optimal labeling L looks like the original image



The Hard Computational Problem

$$\text{data term} \quad \quad \quad \text{prior term}$$
$$E(L) = \sum_p D_P(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

- Now that we have an energy function, the big question is how do we minimize it?
- Exhaustive search is exponential: if n is the number of pixels, there are 2^n possible labelings L



Graph Cuts for Binary Energy Minimization

- Can minimize this energy exactly with graph cuts!

$$E(L) = \sum_p D_p(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

- Build a graph with 2 terminals s, t

Image pixels are nodes in the graph

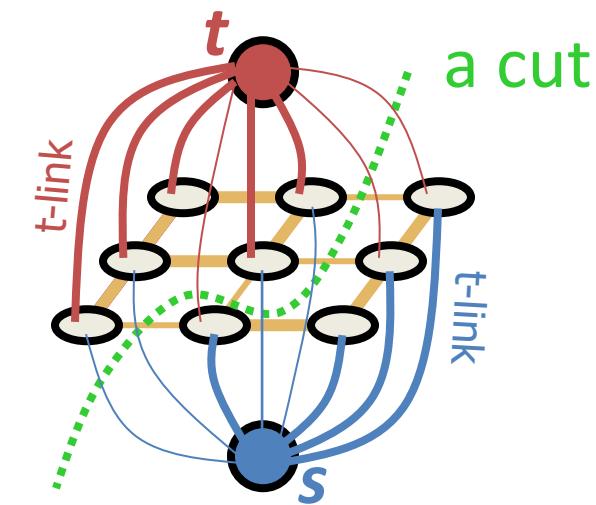
nearby pixels(nodes) connected by an edge, which we call n-link

Terminal s is identified with label 0, and connected by edge called t-link with every image pixel

Terminal t is identified with label 1 and connected by edge called t-link with every image pixel

A cut separates t from s

Each pixel stays connected to either t or s (label 1 or 0)



Graph Cuts for Binary Energy Minimization

$$E(L) = \sum_p D_p(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

■ Build a graph with 2 terminals s, t

Image pixels are nodes in the graph

nearby pixels(nodes) connected by an edge, which called an n-link

Terminal s is identified with label 0, and connected by edge called t-link
with every image pixel

Terminal t is identified with label 1 and connected by edge called t-link with
every image pixel

A cut separates t from s

Each pixel stays connected to either t or s

If pixel p stays connected to terminal s , assign label 0 to p

If pixel p stays connected to terminal t , assign label 1 to p

Cuts correspond to labelings, and with right edge weights cost is same

Thus the minimum cut gives the minimum energy labeling



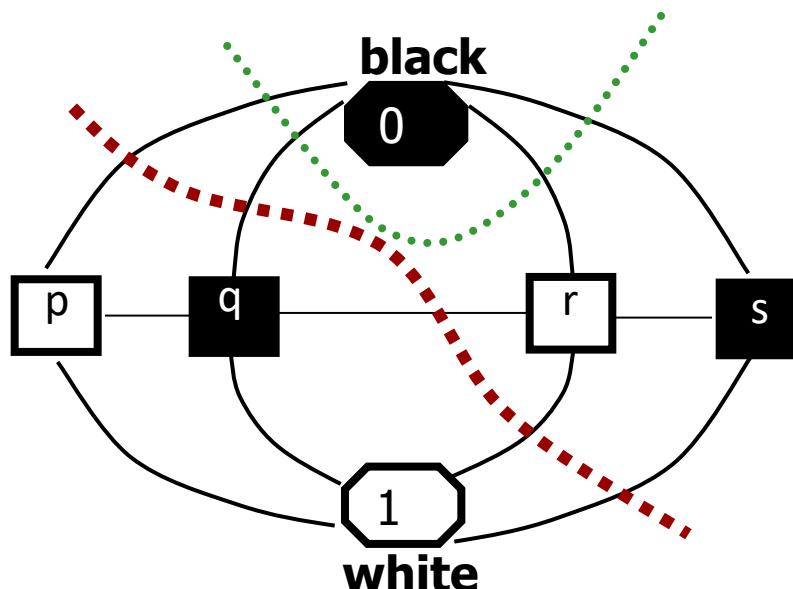
Example

$$E(L) = \sum_p D_p(L_p) + \lambda \sum_{pq \in N} (L_p \neq L_q)$$

- For clarity, let's look at 1 dimensional example but everything works in 2 or higher dimensions. Suppose our image has 4 pixels:



- We build a graph:



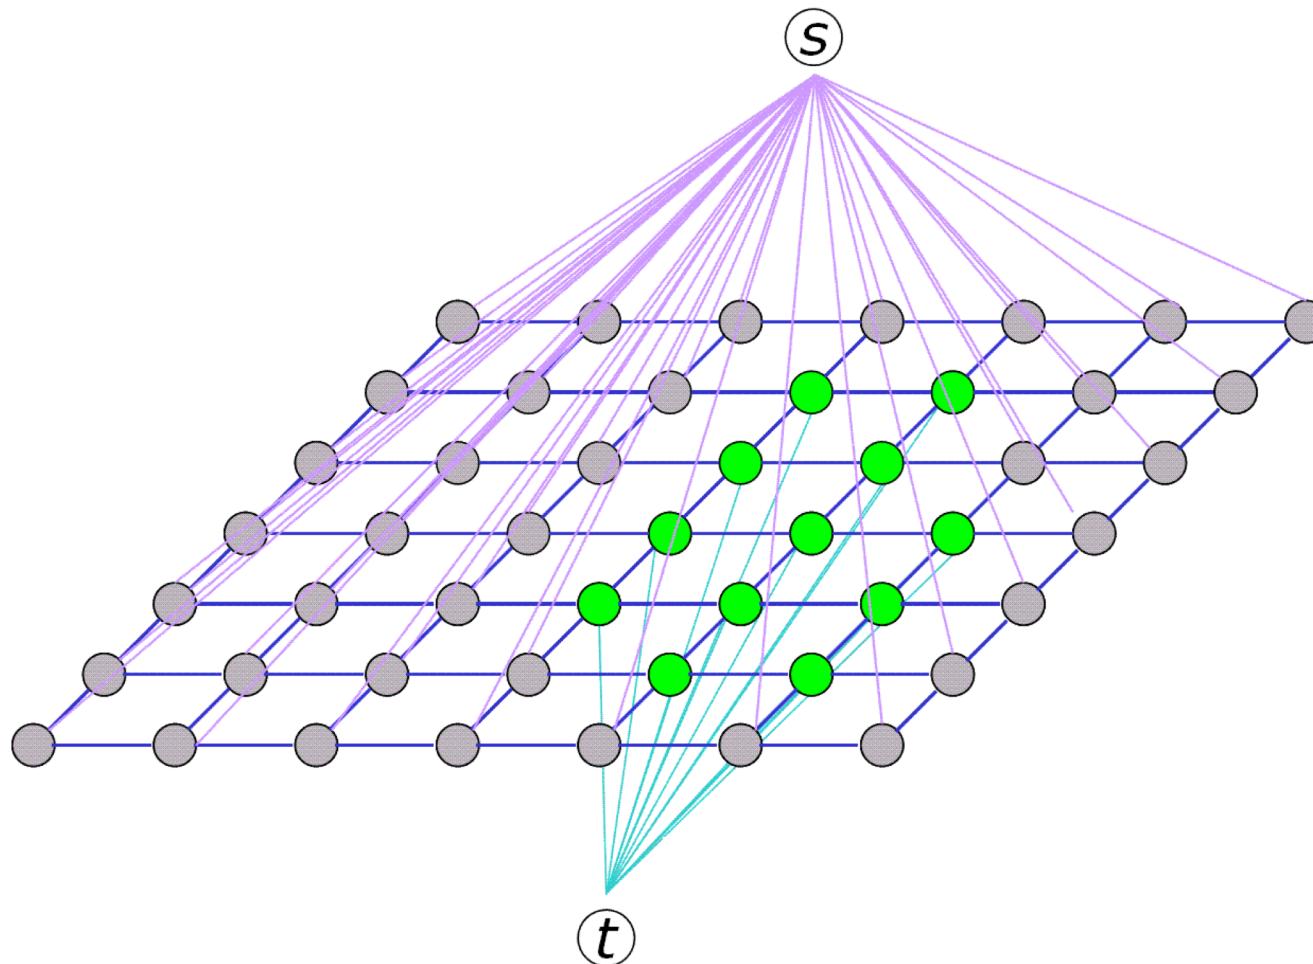
- The cut in red corresponds to labeling



- The cut in green corresponds to labeling



Graph construction



Beyond binary faxes

- All we really need is a cost function
 - Suppose that label 1 means “foreground” and label 0 means “background”
- How do we figure out if a pixel prefers to be in the foreground or background?
 - Predefined intensities: for instance, the foreground object tends to have intensities in the range 50-75
 - So if we observed an intensity in this range at p , $D_p(1)$ is small
 - This is **not** image thresholding!



Better intensity models

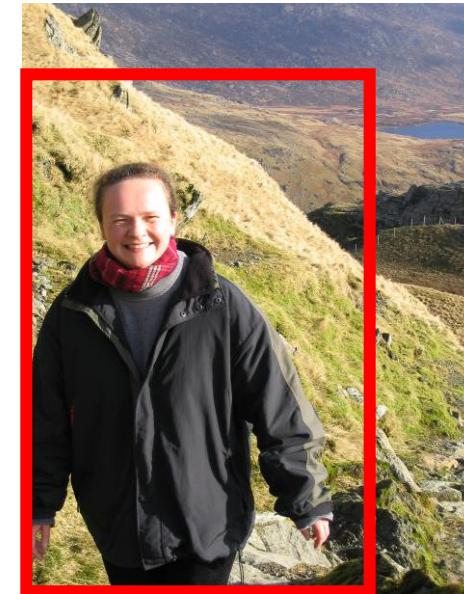
- We can compute the range of intensities dynamically rather than statically
 - Both for foreground and for background
- User marks some pixels as being foreground, and some as background
 - Compute a D_p based on this
 - For instance, $D_p(1)$ is small if p looks like pixels marked as being foreground
- Based on the resulting segmentation, mark additional pixels



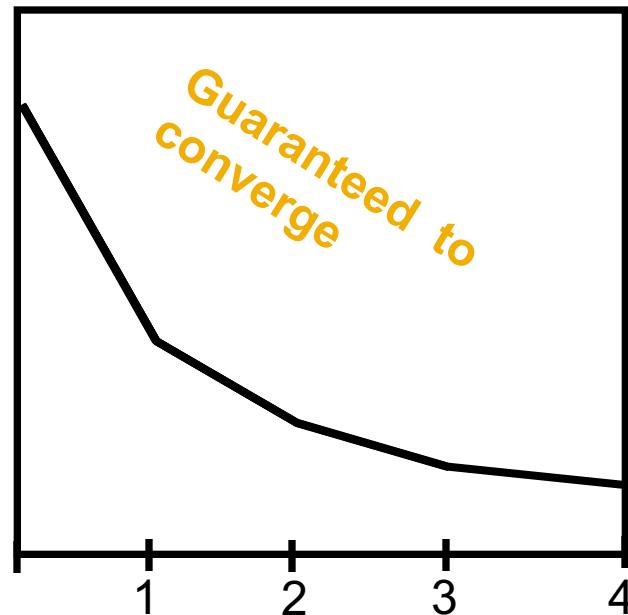
A serious implementation

- This basic segmentation algorithm is now in Microsoft office
 - SIGGRAPH paper on “GrabCut”
- Same basic idea, but simpler UI
 - Assume the background is outside and the foreground is inside a user-supplied box
 - $D_p(1)$ is small if p looks like the pixels inside the box, large if like the pixels outside
 - Create a new segmentation and use it to re-estimate foreground and background

GrabCut



Iterated Graph Cuts



Result

See: <http://www.youtube.com/watch?v=9jNB6fza0nA&feature=related>

Moderately straightforward examples



... GrabCut completes automatically



CORNELL
TECH

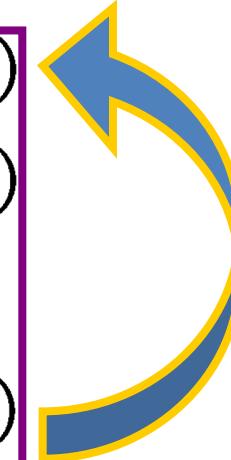
Can this be generalized?

- NP-hard for Potts model
- Two main approaches
 1. Exact solution [Ishikawa 03]
 - Large graph, convex V (arbitrary D)
 - Not the considered the right prior for vision
 2. Approximate solutions
 - Solve a binary labeling problem, repeatedly
 - Expansion move algorithm



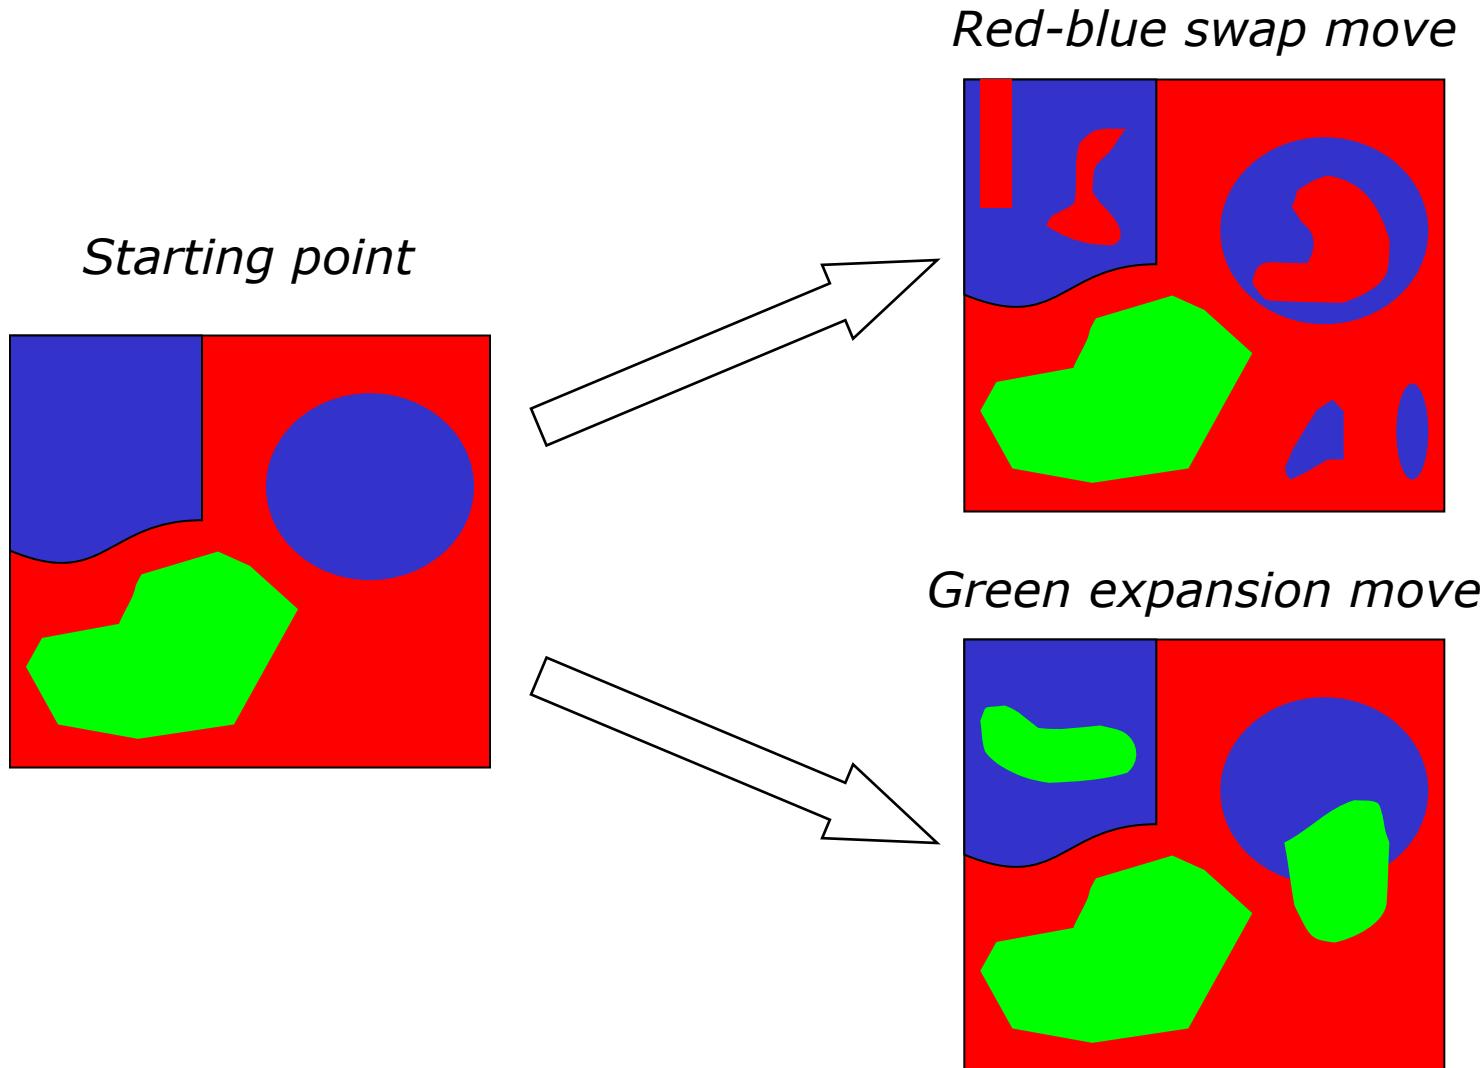
Coordinate descent (local improvement)

$$\begin{aligned}x_1 &:= \arg \min_x E(x, x_2, x_3, \dots, x_n) \\x_2 &:= \arg \min_x E(x_1, x, x_3, \dots, x_n) \\&\vdots \\x_n &:= \arg \min_x E(x_1, x_2, x_3, \dots, x)\end{aligned}$$



- Subproblem: pick a pixel, find the label that minimizes E , repeat
 - Minimize restricted version of E (line search)
 - Computes a local minimum

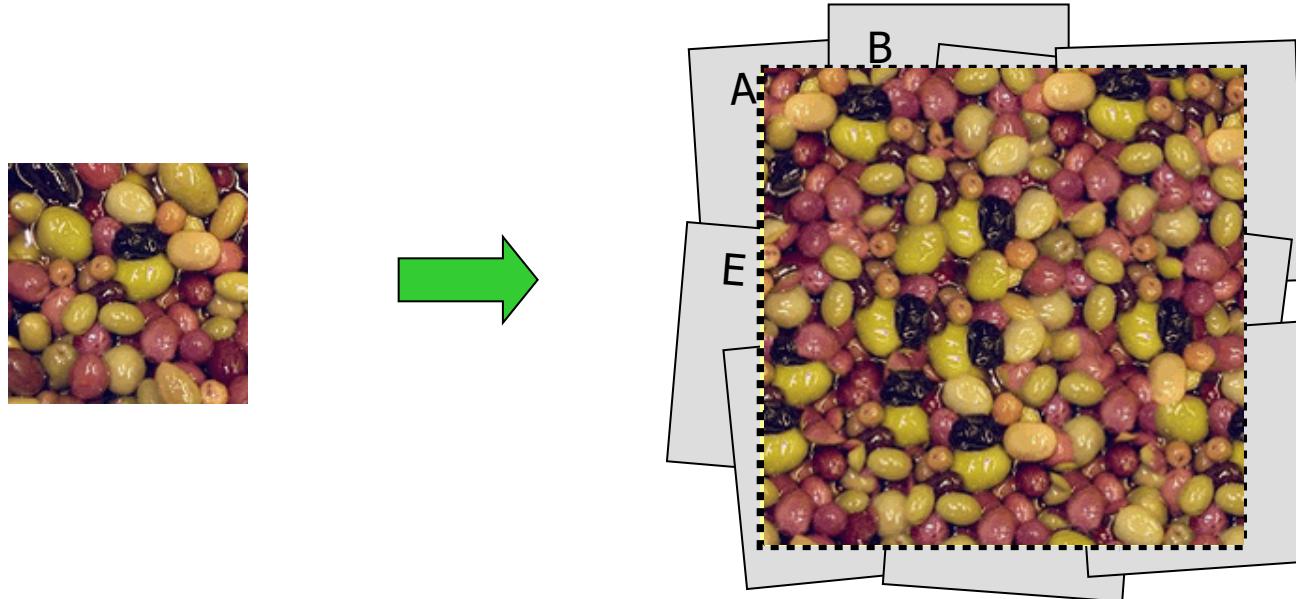
Swap moves and expansion moves



Important properties

- Very efficient in practice
 - Lots of short paths, so roughly linear
- Construction is symmetric (0 vs 1)
- Specific to 2 labels
 - Min cut with >2 labels is NP-hard

Application: texture synthesis

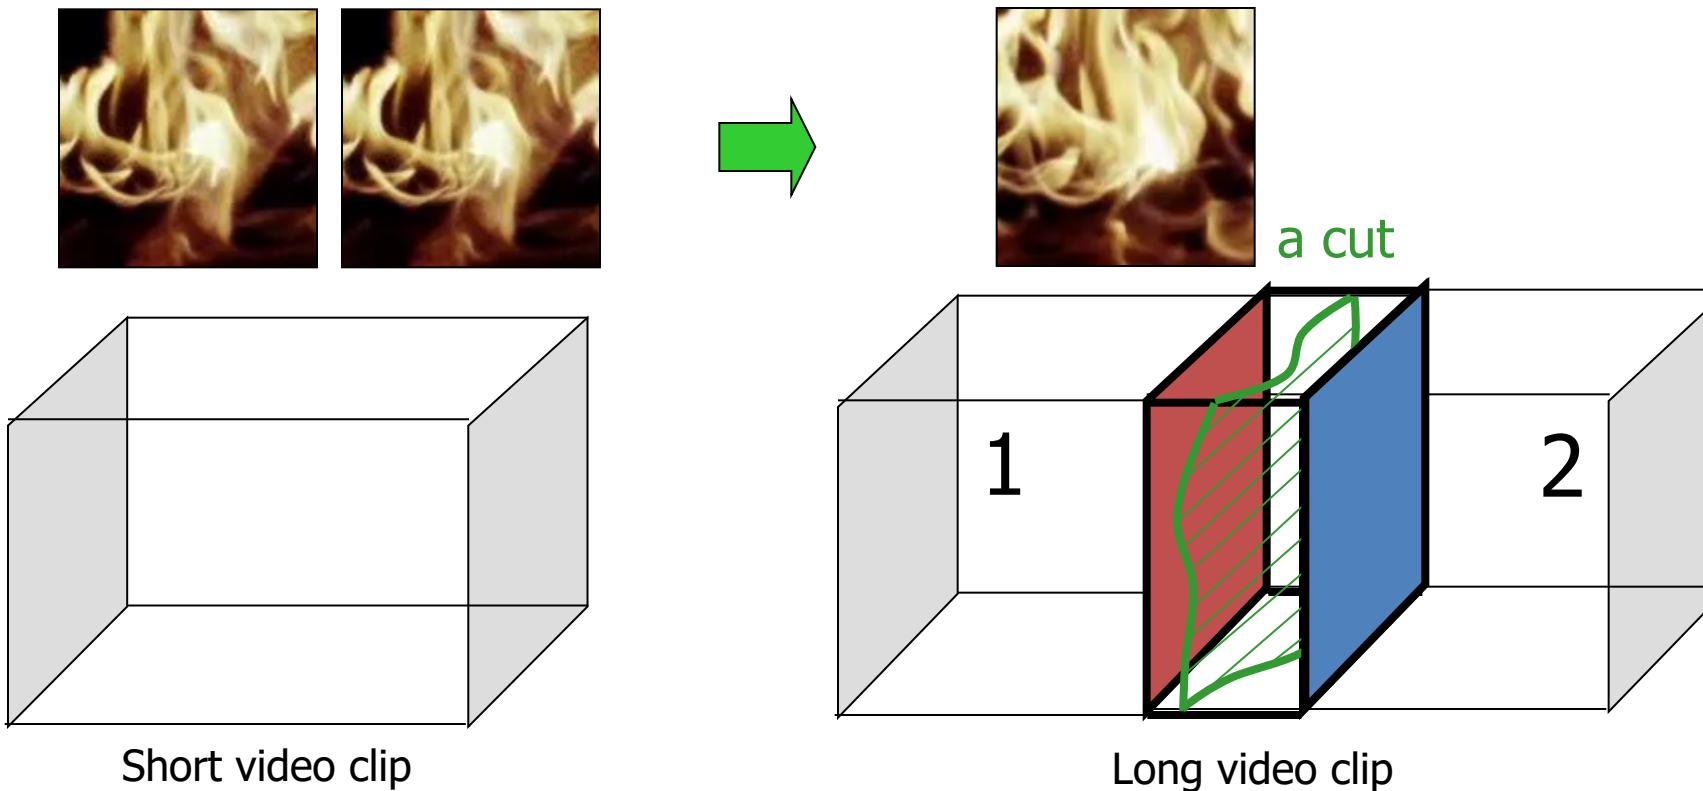


“Graphcut textures” [Kwatra et al 03]



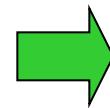
CORNELL
TECH

Graphcuts video textures



Another example

original short clip



synthetic infinite texture



CORNELL
TECH



set of originals

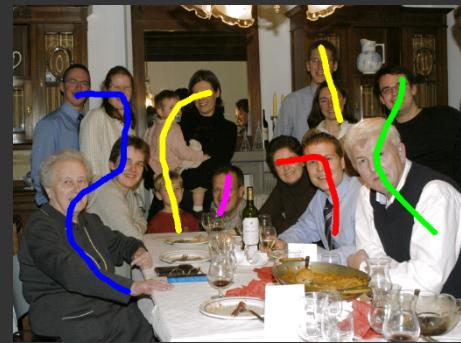


photovintage

Source images



Brush strokes



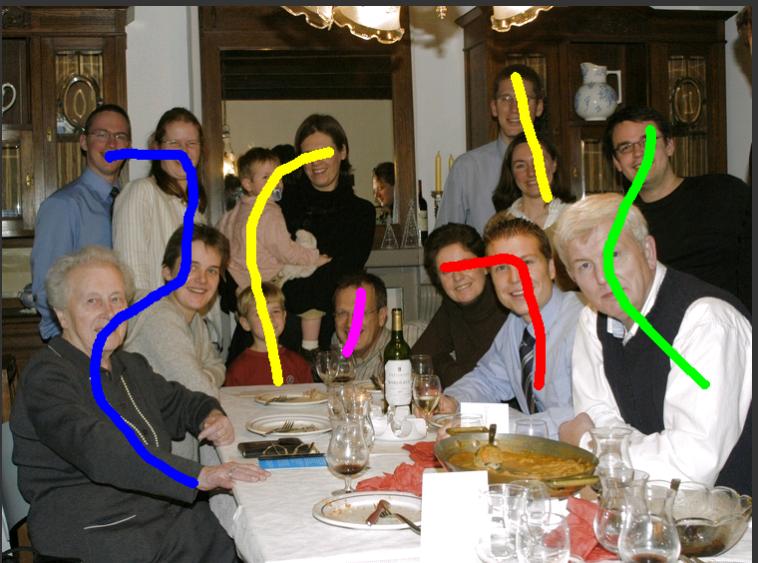
Computed labeling



Composite



Brush strokes



Computed labeling



Image objective

